

Assignment 2 - Logistic Regression Model for Diabetes Prediction

This assignment focuses on building a **Logistic Regression model** to predict diabetes using a dataset ('updated dataset.csv'). The workflow includes data preprocessing, feature encoding, model training, evaluation, and prediction using a sample input.

Code Steps:

1. Import necessary libraries: numpy, pandas, sklearn, pickle.
2. Load dataset and handle missing values using a `safe_fillna()` function.
3. Rename columns for better readability and consistency.
4. Prepare features (X) and target (y). Encode categorical features with one-hot encoding.
5. Split data into training and testing sets (80/20).
6. Standardize features using `StandardScaler`.
7. Train a Logistic Regression model (`max_iter=1000`).
8. Evaluate the model using:
 - Accuracy Score
 - Confusion Matrix
 - Classification Report
9. Save the trained model, scaler, and feature names using pickle.
10. Perform prediction on a new sample input and print whether it's 'Diabetic' or 'Not Diabetic'.

Expected Output:

- Data loaded successfully
- Model training completed
- Accuracy: Displayed as a decimal (e.g., 0.86)
- Confusion Matrix: Printed as a table of true/false positives/negatives
- Classification Report: Displays precision, recall, f1-score
- Model, Scaler & Features saved successfully
- Prediction for sample input: "Diabetic" or "Not Diabetic"

Conclusion:

This assignment demonstrates the complete machine learning pipeline for diabetes prediction. It involves preprocessing data, training a Logistic Regression model, evaluating its performance, and making new predictions. The model, scaler, and feature set are serialized for future use.